


# RF out-of-bag samples

Validation sets for free!


Terence Parr  
MSDS program  
**University of San Francisco**

# RF's have built-in out-of-bag validation set


- RFs have a major advantage over other models: OOB metrics
- Each tree is trained on 63% of data, leaving 37% OOB
- OOB record subsets available to each tree is different
- It's an excellent estimate of the validation error
- Stick with OOB unless default sklearn score() is not suitable
- Not having to process training and validation sets separately is a huge productivity win (assuming significant feature engineering)



# Computing OOB predictions



- Get  $\hat{y}^{(i)}$  by averaging estimates from trees not trained with  $(x^{(i)}, y^{(i)})$ 
  - Image to right; blue is training set, OOB orange
  - Trees from same labeled OOB region of  $x^{(i)}$  used to get  $\hat{y}^{(i)}$
  - Must find all trees not trained on  $x^{(i)}$
  - E.g., compute  $\hat{y}^{(i)}$  for **B** region using Trees 0, 1 but not 2
  - No OOB error estimate is possible for unlabeled regions
- Do not make predictions for OOB and compute error per tree!
- Each tree has high bias, so OOB scores from one tree would be very high
- Average OOB predictions to get  $\hat{y}$  then compute metric on predicted  $\hat{y}$  vector as usual
- Algorithms for regression and classification later



# OOB continued

- OOB error might slightly overestimate test set error. Why?
  - OOB samples are not predicted with all trees in forest whereas test set uses whole forest, which presumably has lower bias/variation [1]
- Some research suggests OOB overestimates error for binary classification <https://journals.plos.org/plosone/article?id=10.1371/journal.pone.0201904>
- OOB metrics don't affect training, just gives metric
- OOB not to be used with time-sensitive data sets. Why not? Validation set for time-sensitive data can't be split randomly

[1] For  $n \ll p$  case, see paper [https://file.scirp.org/Html/9-1240025\\_8072.htm](https://file.scirp.org/Html/9-1240025_8072.htm)



# When OOB error is lower than validation

- Maybe the validation set is drawn from a different distribution than the training set or it's a time-sensitive data set
- Or, the model is overfit to the data in the training set, focusing on relationships that are not relevant to the test set
  - E.g., dropping SalesID transaction ID from training set improved our RF model as SalesID never seen in valid set but predictive in training set
  - E.g., if we add (predictive) feature and OOB is still ok, but the validation set is worse, the validation set is not good
- (Sometimes the validation score is a bit better or worse than the OOB score, due to random fluctuations caused by the inherent randomness of RF construction)


# Extremely randomized trees (Geurts *et al* 2006)

- The variable/value pair is highly sensitive to the training set, and responsible for much of error rate
- “*The optimal cut-point was shown to depend very strongly on the particular learning sample used...this cut-point variance appeared to be responsible for a significant part of the error rates of tree-based methods.*” <https://link.springer.com/article/10.1007/s10994-006-6226-1>
- Geurts wondered if more randomness could reduce variance further
- Pick random split value in  $\min(X[:,j]) \dots \max(X[:,j])$ , ignoring  $y$ !
- Like RF, select  $m \leq p$  variables and choose var/value with lowest loss
- Fits using entire  $X$  training set, not bootstrap and not subsample (trying to reduce bias)
- Our use of just 11 (not  $n$ )  $X$  candidate values in the project is similar (an effort to reduce variance and increase speed)

# OOB regression scoring

**Algorithm:**  $oob\_score_{regr}(RF, X, y)$

**Let**  $oob\_counts[i] = 0 \ \forall \text{ records } i = 1..|X|$  (*Num obs. in all leaves reached by  $X[i]$* )  
**Let**  $oob\_preds[i] = 0 \ \forall \text{ records } i = 1..|X|$  (*Predictions for  $X[i]$  weighted by leaf size*)  
**foreach**  $t \in RF$  **do**  
     $leafsizes = |t.leaf(X[t.oob])|$  (*Num samples in leaf reached by each  $X$* )  
     $oob\_preds[t.oob] += leafsizes \otimes t.predict(X[t.oob])$   
     $oob\_counts[t.oob] += leafsizes$   
**end**  
 $oob\_avg\_preds = \frac{oob\_preds[oob\_counts > 0]}{oob\_counts[oob\_counts > 0]}$   
**return**  $R^2$  score for  $(y[oob\_counts > 0], oob\_avg\_preds)$



Assumes each tree collects OOB sample indexes during fit()

# OOB classification scoring

**Algorithm:**  $oob\_score_{class}(RF, X, y)$

**Let**  $oob\_counts[i] = 0 \ \forall \text{ records } i = 1..|X|$  (Num trees w/predictions for  $X[i]$ )  
 (Create 2D matrix tracking vote counts per class for each  $X[i]$ ):

**Let**  $oob\_preds[i, k] = 0 \ \forall \text{ records } i = 1..|X|, k = 1..|unique(y)|$

**foreach**  $t \in RF$  **do**

$leafsizes = |t.leaf(X[t.oob])|$  (Num samples in leaf reached by each OOB  $X$ )

$tpred = t.predict(X[t.oob])$

$oob\_preds[t.oob, tpred] += leafsizes$  (count weighted class votes)

$oob\_counts[t.oob] += 1$  (track num trees used for each OOB  $X$ )


**end**

**for**  $i$  such that  $oob\_counts[i] > 0$  **do**

$oob\_votes[i] = \arg \max_k oob\_preds[i, k]$

**end**

**return** accuracy of  $y[oob\_counts > 0] = oob\_votes$



Assumes each tree collects OOB sample indexes during fit()