

# Training decision trees

Terence Parr  
MSDS program  
**University of San Francisco**

# Training overview

- Training partitions feature space into rectangular hypervolumes of similar  $X$  records chosen so the associated  $y$  are similar/pure
- Hypervolumes are specified by sequence of *splits* that test a single feature and feature value at a time
- Each split becomes a decision node in decision tree
- Records in an “atomic” hypervolume form a single leaf
- Hypervolume described by conditionals on path from root to leaf

# How to create a decision node

- Given  $(X, y)$  for entire training set or a subregion
- Each split chosen greedily to minimize impurity in subregion  $y$ 's
  - Regressor: variance or MSE
  - Classifier: gini impurity or entropy
- To choose split, exhaustively try each  $(x_j \text{ variable}, x_j \text{ value})$  pair and pick the pair with min weighted average impurity for the two subregions created by that split

# Fitting decision trees

**Algorithm:** *dtreefit*(*X*, *y*, *loss*, *min\_samples\_leaf*)

*if*  $|X| \leq \text{min\_samples\_leaf}$  **then** return Leaf(*y*)

*col*, *split* = *bestsplit*(*X*, *y*, *loss*)

*if* *col* = -1 **then** return Leaf(*y*)    (*No better split?*)

*lchild* = *dtreefit*(*X*[*X*<sub>*col*</sub> ≤ *split*], *y*[*X*<sub>*col*</sub> ≤ *split*], *loss*, *min\_samples\_leaf*)

*rchild* = *dtreefit*(*X*[*X*<sub>*col*</sub> > *split*], *y*[*X*<sub>*col*</sub> > *split*], *loss*, *min\_samples\_leaf*)

**return** *DecisionNode*(*col*, *split*, *lchild*, *rchild*)

subsets    MSE or gini function

Optimization: also check if *y* are all same or very close

Overall fit: pass in full (*X*, *y*) to *dtreefit*() and get back the decision tree

# Best split var/value

**Algorithm:** *bestsplit*( $X, y, loss$ )

$best = (col = -1, split = -1, loss = loss(y))$

**for**  $col = 1..p$  **do**

**foreach**  $split \in X_{col}$  **do**

$yl = y[X_{col} \leq split]$

$yr = y[X_{col} > split]$

**if**  $|yl| < min\_samples\_leaf$  **or**  $|yr| < min\_samples\_leaf$  **then continue**

$l = \frac{|yl| \times loss(yl) + |yr| \times loss(yr)}{|y|}$       (*weighted average of subregion losses*)

**if**  $l = 0$  **then return**  $col, split$

**if**  $l < best[loss]$  **then**  $best = (col, split, l)$

**end**

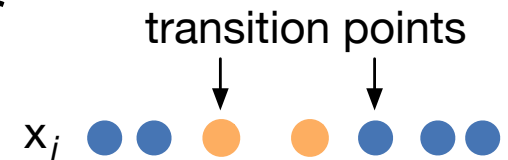
**end**

**return**  $best[col], best[split]$

Should pick midpoint between  
split value and next smallest  $x$

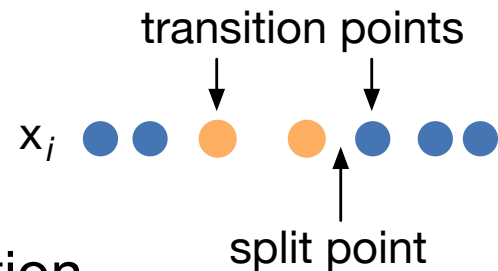
# The usual bestsplit() is inefficient

- It has a nested loop; tries all combinations of  $p$  variables and worst-case  $n$  unique values in each variable at root:  $O(np)$
- Cost of computing loss on all values in subregion each iteration is also expensive
- For classification, can mitigate by sorting by  $i$ th var then we know at a specific  $x$  value, everything to left is less and right is greater; keep track of class counts to left/right
- Reduce computation by focusing on transitions points in  $x$ , effectively focusing on  $\text{unique}(x)$



# Improving generality and efficiency

- Select a subset of values as candidates,  $k$ ; then we reduce  $O(np)$  to  $O(kp)$  for  $k \ll n$  ( $n$  is often huge) (our project  $k=11$ )
- We should really pick split point between two  $x$  values:  $(x^{(i)} + x^{(i-1)})/2$  (if sorted)
- More likely split point is between, not on,  $x$  values, so midpoint is good guess as to underlying distribution
- And, of course, we can reduce tree height with `min_samples_leaf` to restrict complexity



# Decision tree prediction via x subset

**Algorithm:** *bestsplit*( $X, y, loss$ )

$best = (col = -1, split = -1, loss = loss(y))$

**for**  $col = 1..p$  **do**

$candidates =$  randomly pick  $k \ll n$  values from  $X_{col}$

**foreach**  $split \in candidates$  **do**

$yl = y[X \leq split]$

$yr = y[X > split]$

**if**  $|yl| < min\_samples\_leaf$  **or**  $|yr| < min\_samples\_leaf$  **then continue**

$l = \frac{|yl| \times loss(yl) + |yr| \times loss(yr)}{|y|}$  (weighted average of subregion losses)

**if**  $l = 0$  **then return**  $col, split$

**if**  $l < best[loss]$  **then**  $best = (col, split, l)$

**end**

**end**

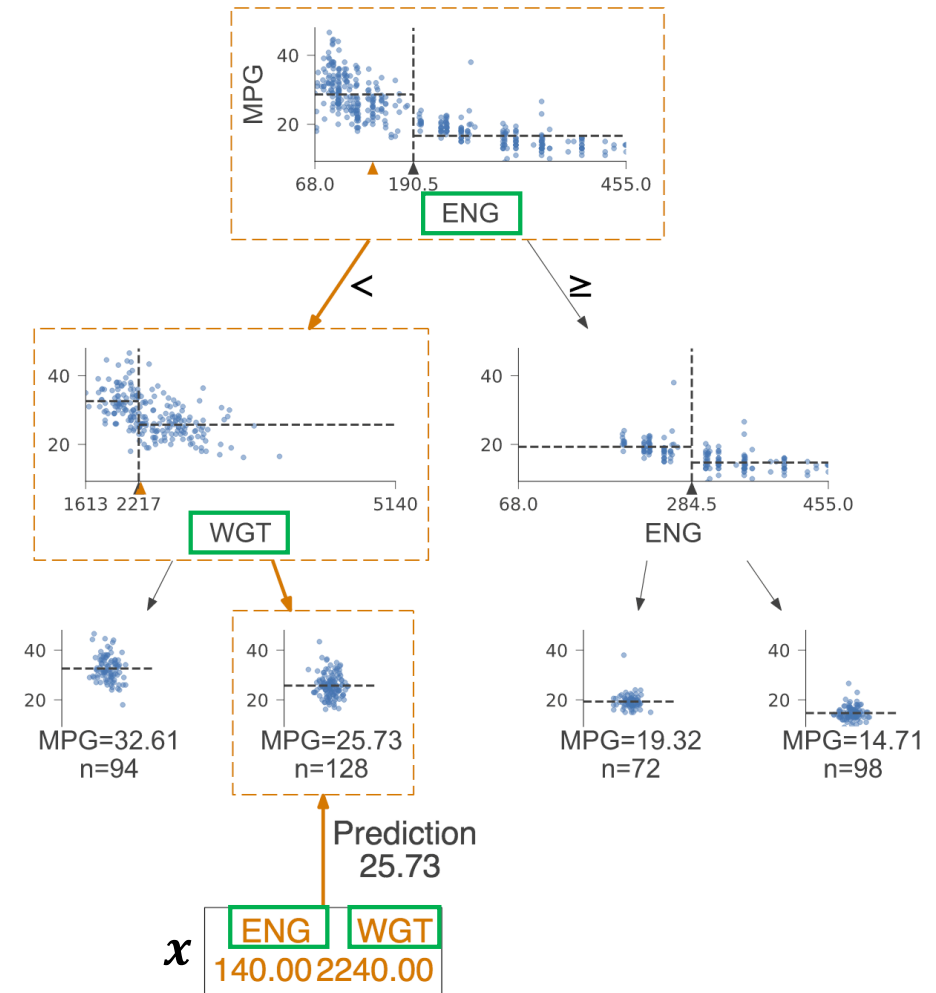
**return**  $best[col], best[split]$

Can even pick just 1 split randomly or in min..max range (see “Extremely random trees”); any small  $k$  value works.



# Prediction

- Start at the root node with test  $x$  and descend through decision nodes to the appropriate leaf; predict leaf mean or mode
- At each decision node, test indicated variable's  $x_j$  value against the split value stored in the decision node



# Prediction algorithm

```
1 Algorithm: predict(node, x)
2   if node is leaf then
3     if classifier then return mode(node.y)
4     return mean(node.y)
5   end
6   if  $x[\text{node.col}] \leq \text{node.split}$  then return predict(node.lchild, x)
7   return predict(node.rchild, x)
```

*y* for samples in leaf

